

Random Forest

Carrie Cheng

2023-04-30

```
library(dplyr)
dat <- read.csv('brfss_final.csv')
outcome <- data.frame(dat$X, dat$MICH, dat$CVDINFR4, dat$CVDCRHD4)
outcome %>% group_by(dat.MICH) %>% summarise(count=n())
```

```
## # A tibble: 2 x 2
##   dat.MICH count
##   <int> <int>
## 1     1  14580
## 2     2  14580
```

```
outcome %>% group_by(dat.CVDINFR4) %>% summarise(count=n())
```

```
## # A tibble: 4 x 2
##   dat.CVDINFR4 count
##   <int> <int>
## 1         1   9188
## 2         2  19802
## 3         7    160
## 4         9    10
```

```
outcome %>% group_by(dat.CVDCRHD4) %>% summarise(count=n())
```

```
## # A tibble: 4 x 2
##   dat.CVDCRHD4 count
##   <int> <int>
## 1         1  9729
## 2         2 18874
## 3         7   550
## 4         9    7
```

```
## remove the ones that responded don't know & not sure in CVDINFR4 & CVDCRHD4
dat <- dat[-which(dat$CVDINFR4 == 7 | dat$CVDINFR4 == 9),]
dat <- dat[-which(dat$CVDCRHD4 == 7 | dat$CVDCRHD4 == 9),]
# remove columns that has only 1 value for all rows
dat <- dat[, -which(names(dat) %in% c("MEDSHEPB", "TOLDCFS", "HAVECFS", "WORKCFS"))]
```

Drop columns with more than 5% data missing, impute the rest using KNN

```

# convert outcome variables
dat$MICHHD <- factor(2-dat$MICHHD)
dat$CVDINFR4 <- factor(2-dat$CVDINFR4)
dat$CVDCRHD4 <- factor(2-dat$CVDCRHD4)
# i believe X is the index column, not needed
# remove weights
dat <- dat[, !colnames(dat) %in% c('X', 'LLCPWT2', 'LLCPWT', 'CLLCPWT', 'STRWT', 'WT2RAKE')]
dat <- dat[, !colnames(dat) %in% c('QSTVER', 'STSTR', 'RAWRAKE')] # remove based on knowledge
threshold <- .05
ncol(dat) # 190

```

```
## [1] 187
```

```

dat <- dat[, colMeans(is.na(dat)) <= threshold]
ncol(dat) # 52 columns left

```

```
## [1] 49
```

```

columns_to_impute <- colnames(dat)[colSums(is.na(dat)) > 0]
columns_to_impute

```

```

## [1] "CPDEMO1B" "VETERAN3" "EMPLOY1" "INCOME3" "DEAF" "BLIND"
## [7] "DECIDE" "DIFFWALK" "DIFFDRES" "DIFFALON" "USENOW3" "METSTAT"
## [13] "URBSTAT" "MSCODE" "DRDXAR3"

```

```
str(dat[,columns_to_impute])
```

```

## 'data.frame': 28433 obs. of 15 variables:
## $ CPDEMO1B: int 1 1 8 1 1 8 8 1 1 2 ...
## $ VETERAN3: int 2 2 2 2 1 2 1 2 2 2 ...
## $ EMPLOY1 : int 8 7 2 7 7 7 7 8 7 7 ...
## $ INCOME3 : int 77 3 99 77 7 99 5 77 5 10 ...
## $ DEAF : int 2 2 2 2 2 2 1 2 2 2 ...
## $ BLIND : int 1 2 2 2 2 2 2 2 2 2 ...
## $ DECIDE : int 1 2 1 2 1 2 2 2 2 2 ...
## $ DIFFWALK: int 1 2 2 2 2 1 1 1 2 2 ...
## $ DIFFDRES: int 2 2 2 2 2 1 2 2 2 2 ...
## $ DIFFALON: int 1 2 2 2 2 1 1 2 2 2 ...
## $ USENOW3 : int 3 3 3 3 3 3 3 3 3 3 ...
## $ METSTAT : int 1 1 1 1 1 2 1 2 1 1 ...
## $ URBSTAT : int 1 1 1 1 1 1 1 1 1 1 ...
## $ MSCODE : int 2 1 3 1 3 2 2 5 2 3 ...
## $ DRDXAR3 : int 1 2 1 1 2 1 1 2 1 1 ...

```

```

complete_columns <- colnames(dat)[colSums(is.na(dat)) == 0 &
!colnames(dat) %in% c('MICHHD', 'CVDINFR4', 'CVDCRHD4')]
for (c in columns_to_impute) {
  col <- dat[[c]]
  scaled <- scale(dat[, complete_columns])
  knn <- knn(

```

```

    train = scaled[!is.na(col), complete_columns],
    test  = scaled[is.na(col), complete_columns],
    cl    = dat[!is.na(col), c]
  )

  dat[is.na(col), c] = knn
}
colSums(is.na(dat))

```

```

## GENHLTH PHYSHLTH MENTHLTH PRIMINSR PERSDOC3 MEDCOST1 CHECKUP1 CVDINFR4
##      0      0      0      0      0      0      0      0      0
## CVDCHRHD4 CVDSTRK3 CHCSCNCR CHCOCNCR CHCCOPD3 ADDEPEV3 CHCKDNY2 DIABETE4
##      0      0      0      0      0      0      0      0      0
## MARITAL  RENTHOM1 NUMHHOL3 CPDEMO1B VETERAN3  EMPLOY1  INCOME3    DEAF
##      0      0      0      0      0      0      0      0      0
##    BLIND    DECIDE DIFFWALK DIFFDRES DIFFALON  USENOW3  QSTLANG  METSTAT
##      0      0      0      0      0      0      0      0      0
## URBSTAT  MSCODE  DUALUSE  TOTINDA  RFHYPE6  CHOLCH3  MICHHD  ASTHMS1
##      0      0      0      0      0      0      0      0      0
## DRDXAR3    RACE    SEX    AGE80  CHLDCNT  EDUCAG  SMOKER3  CURECI1
##      0      0      0      0      0      0      0      0      0
## DROCDY3_
##      0

```

```
library(caret)
```

```
## Loading required package: ggplot2
```

```
## Loading required package: lattice
```

```
library(randomForest)
```

```
## randomForest 4.7-1.1
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##
```

```
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:ggplot2':
```

```
##
```

```
##      margin
```

```
## The following object is masked from 'package:dplyr':
```

```
##
```

```
##      combine
```

```

library(ggplot2)
library(ROCR)
set.seed(263)
train_index <- createDataPartition(dat$MICHHD, p = 0.8, list = FALSE)
train <- dat[train_index, ]
test <- dat[-train_index, ]
train$weights <- ifelse(as.numeric(train$MICHHD) == 1,
                        1/mean(as.numeric(train$MICHHD) == 1),
                        1/(1-mean(as.numeric(train$MICHHD) == 1)))
train$weights <- as.numeric(train$weights)
test$weights <- ifelse(as.numeric(test$MICHHD) == 1,
                       1/mean(as.numeric(test$MICHHD) == 1),
                       1/(1-mean(as.numeric(test$MICHHD) == 1)))
test$weights <- as.numeric(test$weights)
index_weight <- which(names(train) == "weights")

```

Parameter Tuning

Let's tune number of trees `ntrees` and number of features selected to place split `mtry`. In the following, let's use 10-fold cross-validation.

```

## get index of the other two outcomes

index_michd <- which(names(train) == "MICHHD")
index_infr <- which(names(train) == "CVDINFR4")
index_crhd <- which(names(train) == "CVDICRHD4")

```

Tune number of trees

Let's set `mtry = 10`.

```

ntree <- seq(1, 51, by = 20)
accuracy <- sapply(ntree, function(n){
  train(as.factor(MICHHD) ~ ., method = "rf",
        data = train[, -c(index_infr, index_crhd)],
        tuneGrid = data.frame(mtry = 10),
        ntree = n, trControl = trainControl(method = "cv", number = 10))$results$Accuracy
})


qplot(ntree, accuracy)

```

```

## Warning: 'qplot()' was deprecated in ggplot2 3.4.0.
## This warning is displayed once every 8 hours.
## Call 'lifecycle::last_lifecycle_warnings()' to see where this warning was
## generated.

```




```
best_ntree <- ntree[which(accuracy == max(accuracy))]
best_ntree <- min(best_ntree)
print(paste("The best ntree is", best_ntree))
```

```
## [1] "The best ntree is 21"
```

Tune mtry

```
train_rf <- train(as.factor(MICHD) ~ ., method = "rf",
  data = train[, -c(index_infr, index_crhd)],
  tuneGrid = data.frame(mtry = seq(1, 20, by = 5)),
  ntree = best_ntree,
  nodesize = 10, trControl = trainControl(method = "cv", number = 10))

plot(train_rf)
```



```
best_mtry <- train_rf$bestTune
result_cv <- train_rf$results
print(paste("The best mtry is ", best_mtry))
```

```
## [1] "The best mtry is 6"
```

Use the best model to train random forest

The below is the confusion matrix on the test set.

```
rf_best <- randomForest(as.factor(MICHHD) ~.,
                        data = train[, -c(index_infr, index_crhd)],
                        mtry = best_mtry[[1]], ntree = best_ntree, nodesize = 10,
                        weights = train[, index_weight])

pred_test <- predict(rf_best, newdata = test)
cm_test <- confusionMatrix(pred_test, as.factor(test$MICHHD))

cm_test
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    0    1
##           0 2916    0
```

```
##          1      0 2770
##
##          Accuracy : 1
##          95% CI : (0.9994, 1)
##    No Information Rate : 0.5128
##    P-Value [Acc > NIR] : < 2.2e-16
##
##          Kappa : 1
##
##    McNemar's Test P-Value : NA
##
##          Sensitivity : 1.0000
##          Specificity : 1.0000
##    Pos Pred Value : 1.0000
##    Neg Pred Value : 1.0000
##          Prevalence : 0.5128
##    Detection Rate : 0.5128
##    Detection Prevalence : 0.5128
##    Balanced Accuracy : 1.0000
##
##    'Positive' Class : 0
##
```

```
metric_test <- c(cm_test$overall[["Accuracy"]],
                 cm_test$byClass[c("Sensitivity", "Specificity")])

cat(paste("The overall accuracy using the best tuned random forest model is",
          metric_test[1], "\n",
          "Sensitivity is", metric_test[2], "\n",
          "Specificity is", metric_test[3]))
```

```
## The overall accuracy using the best tuned random forest model is 1
## Sensitivity is 1
## Specificity is 1
```

ROC curve

```
library(pROC)
```

```
## Type 'citation("pROC")' for a citation.
```

```
##
```

```
## Attaching package: 'pROC'
```

```
## The following objects are masked from 'package:stats':
```


```
##
```

```
## cov, smooth, var
```

```
roc_rf <- roc(as.numeric(test$MICHHD) ~ as.numeric(pred_test),
              plot = TRUE, print.auc = TRUE)
```

```
## Setting levels: control = 1, case = 2
```

```
## Setting direction: controls < cases
```



```
print(paste("AUC is", as.numeric(roc_rf$auc)))
```

```
## [1] "AUC is 1"
```

Importance Features

```
importance(rf_best)
```

```
##           MeanDecreaseGini
## GENHLTH           360.720462
## PHYSHLTH           43.739617
## MENTHLTH           29.481553
## PRIMINSR           75.467046
## PERSDOC3           52.246191
## MEDCOST1            5.284555
## CHECKUP1           13.956534
## CVDSTRK3           96.536046
```


## CHCSCNCR	12.910459
## CHCOCNCR	12.117497
## CHCCOPD3	64.222906
## ADDEPEV3	9.708542
## CHCKDNY2	30.446637
## DIABETE4	95.214158
## MARITAL	31.696165
## RENTHOM1	10.207899
## NUMHHOL3	11.784564
## CPDEMO1B	18.481098
## VETERAN3	51.136403
## EMPLOY1	64.174316
## INCOME3	43.994866
## DEAF	15.830642
## BLIND	9.694895
## DECIDE	12.299206
## DIFFWALK	80.982118
## DIFFDRES	6.013997
## DIFFALON	12.001085
## USENOW3	7.430666
## QSTLANG	1.421211
## METSTAT	6.495710
## URBSTAT	8.018690
## MSCODE	22.443519
## DUALUSE	6.289627
## TOTINDA	22.608564
## RFHYPE6	306.700294
## CHOLCH3	15.485804
## ASTHMS1	12.242061
## DRDXAR3	15.631956
## RACE	22.757610
## SEX	87.361339
## AGE80	194.670975
## CHLDCNT	15.164496
## EDUCAG	24.949759
## SMOKER3	38.365643
## CURECI1	6.300702
## DROCDY3_	35.322971
## weights	8780.709056

```
varImpPlot(rf_best)
```

rf_best